

twitter_hate_review_src

November 23, 2021

```
[ ]: from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV, StratifiedKFold
from sklearn.utils import class_weight
from collections import Counter
import re
import pandas as pd
from nltk.tokenize import TweetTokenizer
from nltk.corpus import stopwords
from string import punctuation

import warnings
warnings.filterwarnings("ignore")
warnings.filterwarnings(action='ignore', category=DeprecationWarning)
warnings.filterwarnings(action='ignore', category=FutureWarning)
```

```
[ ]: pre_token_replace_vals = [(re.compile(r'@\w+'), ''),
                               (re.compile(r'http\S+'), '')]
stop_words_filename = 'english'
additional_stop_words = ['amp', 'rt']
```

```
[ ]: def read_data_set(data_filename):
    '''function to read dataset and print some information about the dataset'''
    # read csv file into dataframe
    data_df = pd.read_csv(
        data_filename, delimiter=",", encoding="utf-8")

    # print the info of twitter data frame
    print(data_df.head())
    print(data_df.shape)
    print(data_df.columns)
    print(data_df.isnull().sum())
```

```
return data_df
```

```
[ ]: def get_stop_words(stop_words_filename, additional_stop_words):  
    '''function create english language stop words and add additional required_  
    ↪terms'''  
    stop_words = stopwords.words(stop_words_filename) # get stop words for_  
    ↪given language  
    stop_words.extend(additional_stop_words) # add additional stop words  
  
    return stop_words
```

```
[ ]: def pre_token_cleanup(text, replace_vals):  
    '''function to pre-process the tweets'''  
    text = text.lower() # convert to lower case  
  
    # text = replace_with(text, [('&', 'and'), ('>', '>'), ('<', '  
    ↪<')])  
    for replace in replace_vals:  
        text = re.sub(replace[0], replace[1], text)  
  
    text = text.strip() # remove leading and trailing whitespace  
    return text
```

```
[ ]: def remove_stop_words(text, stop_words):  
    '''function to remove stop words for given text'''  
    text = [word for word in text if word not in stop_words]  
  
    return text
```

```
[ ]: def remove_hashtags(text, tokenizer):  
    '''function to remove hashtags from the tweets'''  
    pattern = r'#'  
    joined_text = ' '.join(text)  
    cleaned_text = re.sub(pattern, '', joined_text)  
  
    return tokenizer.tokenize(cleaned_text)
```

```
[ ]: def remove_length_one_words(text):  
    '''function to remove length one words'''  
    text = [word for word in text if len(word) > 1]  
  
    return text
```

```
[ ]: def remove_punctuation(text):  
    '''function to remove punctuation from the tweets'''  
    text = [word for word in text if word not in punctuation]
```

```
return text
```

```
[ ]: def remove_nonalpha(text):  
    '''function to remove non alpha characters'''  
    text = [word for word in text if word.isalpha()]  
  
    return text
```

```
[ ]: def read_twitter_data_set():  
    '''function to read the train dataset and print some information about the_  
    ↪dataset and remove unnecessary columns'''  
    # read csv file into dataframe  
    twitter_hate_df = read_data_set('../data/twitter_hate.csv')  
    twitter_hate_df.drop(["id"], axis=1, inplace=True) # remove id column  
  
    return twitter_hate_df
```

```
[ ]: def clean_tweets(twitter_hate_df):  
    '''function to cleanup teh tweets to prepare for training'''  
    # Normalize the casing, remove user handles and URLs.  
    twitter_hate_df["tweet"] = twitter_hate_df["tweet"].apply(  
        pre_token_cleanup, args=(pre_token_replace_vals,))  
    print(twitter_hate_df.head())  
  
    # Using TweetTokenizer from NLTK, tokenize the tweets into individual terms.  
    tokenizer = TweetTokenizer()  
    twitter_hate_df["tweet"] = twitter_hate_df["tweet"].apply(  
        tokenizer.tokenize)  
    print(twitter_hate_df.head())  
  
    # Remove stop words and redundant terms like 'amp', 'rt', etc..  
    stop_words = get_stop_words(  
        stop_words_filename, additional_stop_words)  
    twitter_hate_df["tweet"] = twitter_hate_df["tweet"].apply(  
        remove_stop_words, args=(stop_words,))  
    print(twitter_hate_df.head())  
  
    # Remove '#' symbols from the tweet while retaining the term.  
    twitter_hate_df["tweet"] = twitter_hate_df["tweet"].apply(  
        remove_hashtags, args=(tokenizer,))  
    print(twitter_hate_df.head())  
  
    # Removing tweets with a length of 1.  
    twitter_hate_df["tweet"] = twitter_hate_df["tweet"].apply(  
        remove_length_one_words)  
    print(twitter_hate_df.head())
```

```

# Removing punctuation.
twitter_hate_df["tweet"] = twitter_hate_df["tweet"].apply(
    remove_punctuation)
print(twitter_hate_df.head())

# Removing non-alpha characters.
twitter_hate_df["tweet"] = twitter_hate_df["tweet"].apply(
    remove_nonalpha)
print(twitter_hate_df.head())

return twitter_hate_df

```

```

[ ]: def find_top_terms_in_tweets(twitter_hate_df):
    '''function to find top 10 terms in tweets'''
    # find top 10 terms in tweets
    top_terms = Counter()
    for tweet in twitter_hate_df["tweet"]:
        top_terms.update(tweet)

    print(top_terms.most_common(10))

```

```

[ ]: def format_for_training(twitter_hate_df):
    '''function to format the data for training'''
    # join the tokens back to string
    twitter_hate_df["tweet"] = twitter_hate_df["tweet"].apply(
        lambda x: ' '.join(x))
    print(twitter_hate_df.head())

    X = twitter_hate_df["tweet"]
    y = twitter_hate_df["label"]

    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42)
    print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)

    return X_train, X_test, y_train, y_test

```

```

[ ]: def transform_to_feature_vector(X_train, X_test):
    '''function to transform the data to feature vector'''
    # Create the TF-IDF vectorizer with the top 5000 terms.
    vectorizer = TfidfVectorizer(max_features=5000)
    X_train_vectorized = vectorizer.fit_transform(X_train)
    X_test_vectorized = vectorizer.transform(X_test)

    print(X_train_vectorized.shape, X_test_vectorized.shape)

    return X_train_vectorized, X_test_vectorized, vectorizer

```

```
[ ]: def train_model(X_train_vectorized, y_train):
    '''function to train the model'''
    # Train the model using the training data.
    model = LogisticRegression()
    model.fit(X_train_vectorized, y_train)

    return model
```

```
[ ]: def make_predictions(model, X_train_vectorized, X_test_vectorized):
    '''function to make predictions'''
    # Make predictions on the train data.
    y_train_pred = model.predict(X_train_vectorized)
    y_test_pred = model.predict(X_test_vectorized)

    return y_train_pred, y_test_pred
```

```
[ ]: def evaluate_model(y_train_pred, y_test_pred, y_train, y_test):
    '''function to evaluate the model'''
    # Evaluate the model using the training data.
    print("Accuracy on training data:", accuracy_score(y_train_pred, y_train))
    print("Accuracy on test data:", accuracy_score(y_test_pred, y_test))

    print(classification_report(y_train, y_train_pred))

    return None
```

```
[ ]: def adjusting_class_imbalance_for_balance_model(X_train_vectorized, y_train):
    '''function to adjust class imbalance'''
    # Adjust class imbalance.
    balanced_model = LogisticRegression(class_weight='balanced')
    balanced_model.fit(X_train_vectorized, y_train)

    return balanced_model
```

```
[ ]: def make_predictions_with_balanced_model(balanced_model, X_train_vectorized, X_test_vectorized):
    '''function to make predictions with balanced model'''
    # Make predictions on the train data.
    y_train_pred = balanced_model.predict(X_train_vectorized)
    y_test_pred = balanced_model.predict(X_test_vectorized)

    return y_train_pred, y_test_pred
```

```
[ ]: def evaluate_model_with_balanced_model(y_train_pred, y_test_pred, y_train, y_test):
    '''function to evaluate the model with balanced model'''
    # Evaluate the model using the training data.
```

```

print("Accuracy on training data:", accuracy_score(y_train_pred, y_train))
print("Accuracy on test data:", accuracy_score(y_test_pred, y_test))

print(classification_report(y_train, y_train_pred))

return None

```

```

[ ]: def get_parameters_for_grid_search():
    '''function to get parameters for grid search'''
    # Get the parameters for grid search.
    # parameters = {'C': [0.001, 0.01, 0.1, 1, 10, 100, 1000],
    #               'penalty': ['l1', 'l2']}

    parameters = {'C': [0.01, 0.1, 1, 10, 100],
                  'penalty': ['l1', 'l2']}
    return parameters

```

```

[ ]: def get_grid_search_model(X_train_vectorized, y_train, parameters):
    '''function to grid search model'''
    # Grid search model.
    grid_search = GridSearchCV(estimator=LogisticRegression(class_weight='balanced',
    ↪max_iter=1000),
                                param_grid=parameters,
                                cv=StratifiedKFold(4),
                                n_jobs=-1,
                                verbose=1,
                                scoring='recall')

    grid_search.fit(X_train_vectorized, y_train)

    print(grid_search.best_params_)
    # print(grid_search.best_score_)
    print(grid_search.best_estimator_)
    # print(grid_search.cv_results_)

    return grid_search

```

```

[ ]: def make_predictions_with_grid_search_model(grid_search_model,
    ↪X_train_vectorized, X_test_vectorized):
    '''function to make predictions with grid search model'''
    # Make predictions on the train data.
    y_train_pred = grid_search_model.predict(X_train_vectorized)
    y_test_pred = grid_search_model.predict(X_test_vectorized)

    return y_train_pred, y_test_pred

```

```

[ ]: twitter_hate_df = read_twitter_data_set()

[ ]: twitter_hate_df = clean_tweets(twitter_hate_df)

[ ]: find_top_terms_in_tweets(twitter_hate_df)

[ ]: X_train, X_test, y_train, y_test = format_for_training(twitter_hate_df)
    X_train_vectorized, X_test_vectorized, vectorizer =
    ↪transform_to_feature_vector(X_train, X_test)

[ ]: model = train_model(X_train_vectorized, y_train)

[ ]: y_train_pred, y_test_pred = make_predictions(model, X_train_vectorized,
    ↪X_test_vectorized)

[ ]: evaluate_model(y_train_pred, y_test_pred, y_train, y_test)

[ ]: balanced_model =
    ↪adjusting_class_imbalance_for_balance_model(X_train_vectorized, y_train)

[ ]: y_train_pred, y_test_pred =
    ↪make_predictions_with_balanced_model(balanced_model, X_train_vectorized,
    ↪X_test_vectorized)

[ ]: evaluate_model_with_balanced_model(y_train_pred, y_test_pred, y_train, y_test)

[ ]: grid_search_model = get_grid_search_model(X_train_vectorized, y_train,
    ↪get_parameters_for_grid_search())

[ ]: y_train_pred, y_test_pred =
    ↪make_predictions_with_grid_search_model(grid_search_model,
    ↪X_train_vectorized, X_test_vectorized)

[ ]: print(classification_report(y_test, y_test_pred))

[ ]:

```